

Problem - 1: Perform a classification task with knn from scratch.

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

print("\n")
print("1. Load the Dataset")
print("\n")

data = pd.read_csv("/content/drive/MyDrive/Untitled folder/diabetes_.csv")

print(data.head())
print(data.info())
print(data.describe())
print(data.isnull().sum())

print("\n")
print("2. Handle Missing Data")
print("\n")

for column in data.columns:
    if data[column].isnull().sum() > 0:
        data[column].fillna(data[column].mean(), inplace=True)

print(data.isnull().sum())

print("\n")
print("3. Feature Engineering")
print("\n")

X = data.drop(columns=["Outcome"]).values
y = data["Outcome"].values

def train_test_split_scratch(X, y, test_size=0.3, random_seed=42):
    np.random.seed(random_seed)
    indices = np.arange(len(X))
    np.random.shuffle(indices)

    split = int(len(X) * test_size)
    test_idx = indices[:split]
    train_idx = indices[split:]

    return X[train_idx], X[test_idx], y[train_idx], y[test_idx]

X_train, X_test, y_train, y_test = train_test_split_scratch(X, y)

print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)

print("\n")
print("4. Implement KNN")
print("\n")

def euclidean_distance(p1, p2):
    if p1.shape != p2.shape:
        raise ValueError("Dimension mismatch")
    return np.sqrt(np.sum((p1 - p2) ** 2))

def knn_predict_single(query, X_train, y_train, k):
    distances = []

    for i in range(len(X_train)):
        dist = euclidean_distance(query, X_train[i])
        distances.append(dist)
```

```
nearest_indices = np.argsort(distances)[:k]
nearest_labels = y_train[nearest_indices]

prediction = np.bincount(nearest_labels).argmax()
return prediction

def knn_predict(X_test, X_train, y_train, k):
    predictions = []
    for x in X_test:
        predictions.append(knn_predict_single(x, X_train, y_train, k))
    return np.array(predictions)

def compute_accuracy(y_true, y_pred):
    return (np.sum(y_true == y_pred) / len(y_true)) * 100

predictions = knn_predict(X_test, X_train, y_train, k=3)
accuracy = compute_accuracy(y_test, predictions)

print(f"Accuracy of KNN (k=3) on Unscaled Data: {accuracy:.2f}%")
```

	768.000000	768.000000	768.000000	768.000000	768.000000
count	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479
std	3.369578	31.972618	19.355807	15.952218	115.244002
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000
75%	6.000000	140.250000	80.000000	32.000000	127.250000
max	17.000000	199.000000	122.000000	99.000000	846.000000

	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000
mean	31.992578	0.471876	33.240885	0.348958
std	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.078000	21.000000	0.000000
25%	27.300000	0.243750	24.000000	0.000000
50%	32.000000	0.372500	29.000000	0.000000
75%	36.600000	0.626250	41.000000	1.000000
max	67.100000	2.420000	81.000000	1.000000
Pregnancies	0			
Glucose	0			
BloodPressure	0			

Problem - 2 - Experimentation:

```

def min_max_scaling(X):
    X_min = X.min(axis=0)
    X_max = X.max(axis=0)
    return (X - X_min) / (X_max - X_min)

X_scaled = min_max_scaling(X)

X_train, X_test, y_train, y_test = train_test_split_scratch(X, y)

pred_unscaled = knn_predict(X_test, X_train, y_train, k=3)
accuracy_unscaled = compute_accuracy(y_test, pred_unscaled)

X_train_s, X_test_s, y_train_s, y_test_s = train_test_split_scratch(X_scaled, y)

pred_scaled = knn_predict(X_test_s, X_train_s, y_train_s, k=3)
accuracy_scaled = compute_accuracy(y_test_s, pred_scaled)

print(f"Accuracy on Unscaled Data (k=3): {accuracy_unscaled:.2f}%")
print(f"Accuracy on Scaled Data (k=3): {accuracy_scaled:.2f}%")

```

Accuracy on Unscaled Data (k=3): 67.39%
 Accuracy on Scaled Data (k=3): 69.13%

Problem - 3 - Experimentation with k:

```

import time

def evaluate_knn_k_values(X_train, y_train, X_test, y_test, k_values):
    accuracies = []
    times = []

    for k in k_values:
        start_time = time.time()

        predictions = knn_predict(X_test, X_train, y_train, k)

        end_time = time.time()

        accuracy = compute_accuracy(y_test, predictions)
        time_taken = end_time - start_time

        accuracies.append(accuracy)
        times.append(time_taken)

    print(f"k={k} | Accuracy={accuracy:.2f}% | Time={time_taken:.6f}s")

    return accuracies, times

k_values = range(1, 16)

print("---- Original Dataset Results ----")
acc_original, time_original = evaluate_knn_k_values(
    X_train, y_train, X_test, y_test, k_values
)

print("\n---- Scaled Dataset Results ----")
acc_scaled, time_scaled = evaluate_knn_k_values(
    X_train_s, y_train_s, X_test_s, y_test_s, k_values
)

```

```
import matplotlib.pyplot as plt

plt.figure()
plt.plot(k_values, acc_original, marker='o', label="Original Data")
plt.plot(k_values, acc_scaled, marker='o', label="Scaled Data")
plt.xlabel("k (Number of Neighbors)")
plt.ylabel("Accuracy (%)")
plt.title("k vs Accuracy")
plt.legend()
plt.grid()
plt.show()

plt.figure()
plt.plot(k_values, time_original, marker='o', label="Original Data")
plt.plot(k_values, time_scaled, marker='o', label="Scaled Data")
plt.xlabel("k (Number of Neighbors)")
plt.ylabel("Prediction Time (seconds)")
plt.title("k vs Prediction Time")
plt.legend()
plt.grid()
plt.show()

best_k_original = k_values[acc_original.index(max(acc_original))]
best_k_scaled = k_values[acc_scaled.index(max(acc_scaled))]

print("---- Optimal k Values ----")
print(f"Best k (Original Data): {best_k_original}")
print(f"Best k (Scaled Data): {best_k_scaled}")
```

----- Original Dataset Results -----

k=1	Accuracy=68.70%	Time=0.939347s
k=2	Accuracy=72.61%	Time=1.522459s
k=3	Accuracy=67.39%	Time=1.515273s
k=4	Accuracy=72.17%	Time=1.300493s
k=5	Accuracy=68.70%	Time=0.908821s
k=6	Accuracy=70.00%	Time=0.947219s
k=7	Accuracy=69.13%	Time=0.911161s
k=8	Accuracy=71.30%	Time=0.862743s
k=9	Accuracy=70.00%	Time=0.916247s
k=10	Accuracy=71.74%	Time=2.041885s
k=11	Accuracy=73.04%	Time=0.984937s
k=12	Accuracy=73.91%	Time=0.872694s
k=13	Accuracy=74.35%	Time=0.910316s
k=14	Accuracy=72.61%	Time=1.294807s
k=15	Accuracy=73.91%	Time=1.554614s

----- Scaled Dataset Results -----

k=1	Accuracy=67.39%	Time=1.528401s
k=2	Accuracy=67.83%	Time=0.966769s
k=3	Accuracy=69.13%	Time=0.888378s
k=4	Accuracy=69.57%	Time=0.908580s
k=5	Accuracy=67.83%	Time=0.881157s
k=6	Accuracy=66.52%	Time=0.885844s
k=7	Accuracy=71.30%	Time=0.877103s
k=8	Accuracy=69.57%	Time=0.894743s
k=9	Accuracy=72.17%	Time=0.900736s
k=10	Accuracy=70.43%	Time=0.898081s
k=11	Accuracy=70.00%	Time=0.919598s
k=12	Accuracy=68.70%	Time=0.912148s
k=13	Accuracy=69.13%	Time=1.404951s
k=14	Accuracy=68.70%	Time=1.497566s
k=15	Accuracy=70.87%	Time=1.508878s

